

Emotion Detection Using Transformer Model With DeepLearning

Divyang Bharatbhai Joshi

Department of Computer Science and Engineering
Parul University, Vadodara, Gujarat
divyangjoshiprofessional@gmail.com

Apoorva Ashokbhai Dhokai

Department of Computer Science and Engineering
Parul University, Vadodara, Gujarat
apoorvadhokai716@gmail.com

Nandini Pavankumar Agrawal

Department of Computer Science and Engineering
Parul University, Vadodara, Gujarat
nandiniagrawal139@gmail.com

Asst. Prof. Pankaj Agrawal

Department of Computer Science & Engineering
Parul University, Vadodara, Gujarat
pankaj.agrawal23906@paruluniversity.ac.in

Amitkumar Rajpoot

Department of Computer Science and Engineering
Parul University, Vadodara, Gujarat
amitkumar.rajpoot33412@paruluniversity.ac.in

Mr. Avinash Sood

CEO, Binary Qubit
Gurugram, Haryana
avikumar.research@gmail.com

Abstract — Emotion detection is essential in NLP for classifying emotions from text. Traditional models like SVMs and RNNs struggle with contextual understanding, while Transformer-based models such as BERT and RoBERT offer significant improvements. This study proposes a Transformer-based deep learning approach for emotion classification into six categories. The dataset undergoes preprocessing and label encoding for better model efficiency. Our model achieves over 90% accuracy, surpassing previous deep learning methods (70-85% accuracy). Results show that Transformers capture semantic nuances better than traditional models. This approach is valuable for sentiment analysis, mental health monitoring, and human-computer interaction. Our findings highlight the superiority of Transformer models in emotion detection tasks.

Keywords — Emotion Detection, Natural Language Processing (NLP), Sentiment Analysis, BERT, RoBERT, Text Classification, Contextual Embeddings, Self-Attention Mechanism, Mental Health Monitoring, Human-Computer Interaction.

I. INTRODUCTION

Emotion detection is a vital task in Natural Language Processing (NLP) that aims to classify emotions based on textual data. With the rapid growth of digital communication, understanding human emotions from text has become essential for applications in opinion mining, mental health monitoring, human-computer interaction, and customer sentiment analysis. Traditional sentiment analysis techniques, such as lexicon-based methods and machine learning models, have shown promising results but struggle with capturing complex linguistic patterns and contextual dependencies. The emergence of Transformer-based models has revolutionized emotion detection, providing state-of-the-art performance by leveraging self-attention mechanisms and deep contextual embeddings.

II. Background and Evolution of Sentiment Analysis

Early sentiment analysis models relied heavily on rule-based approaches and lexicon-based methods (Pang & Lee, 2008). These methods utilized predefined sentiment dictionaries and syntactic structures to determine the polarity of a given text. However, they faced challenges in handling sarcasm, context variations, and domain-specific language.

Later, statistical learning approaches were introduced, where researchers employed Naïve Bayes, Support Vector Machines (SVMs), and Decision Trees for opinion mining (Hu & Liu, 2004). These models improved classification accuracy but still lacked the ability to capture contextual meaning in complex sentences.

With advancements in deep learning, models such as Long Short-Term Memory (LSTM) networks and Convolutional Neural Networks (CNNs) demonstrated better performance in sentiment analysis. However, these models struggled with long-range dependencies and interpretability in textual data.

III. The Rise of Transformer-Based Models

The Transformer architecture, introduced through BERT (Bidirectional Encoder Representations from Transformers), addressed many of the limitations faced by traditional deep learning models. Unlike previous methods, Transformers rely on self-attention mechanisms to understand the context of words in relation to one another (Cambria & Hussain, 2012). This allows them to capture nuanced emotions, sarcasm, and polysemy, significantly improving emotion detection accuracy.

Moreover, Transformer-based models such as RoBERT, XLNet, and GPT have demonstrated superior performance in

concept-level sentiment analysis, where emotion detection is performed based on deeper semantic understanding rather than surface-level word associations (Poria et al., 2017). These models leverage **large-scale pretraining** and **fine-tuning strategies** to adapt to various domains, making them **highly effective in real-world sentiment classification tasks**.

IV. Multimodal Approaches in Emotion Detection

Recent research in **affective computing** has explored **multimodal fusion techniques**, combining textual, visual, and auditory data to enhance sentiment analysis. Poria et al. (2017) reviewed **multimodal sentiment analysis**, highlighting how fusing multiple data sources improves the robustness of emotion classification. By integrating **Transformer-based text models with speech and facial expression analysis**, researchers have achieved **more accurate and holistic emotion detection systems**.

V. Challenges and Future Directions

Despite the advancements brought by Transformer-based models, several challenges remain in **emotion detection**:

- **Data Scarcity**: High-quality emotion-labeled datasets are limited, requiring models to generalize across different domains.
- **Explainability**: Transformer models function as black-box systems, making it difficult to interpret their decision-making process.
- **Computational Costs**: Training and deploying large-scale Transformer models demand significant computational resources.

I. LITERATURE REVIEW

Bo Pang and Lillian Lee[1] provided a comprehensive overview of sentiment analysis and opinion mining techniques. Their survey covers a wide range of methods, including machine learning algorithms, natural language processing (NLP) approaches, and lexicon-based techniques. The paper discusses applications in various domains, such as **social media analysis, product reviews, and recommendation systems**. The authors emphasize the **challenges in opinion mining, including handling sarcasm, ambiguity, and multilingual data**. This work is foundational for understanding the evolution of sentiment analysis techniques and their practical applications.

Minqing Hu and Bing Liu's[2] work introduced an effective method for extracting product features from customer reviews and determining their sentiment orientations. The proposed technique applies data mining approaches to identify product attributes, analyze their frequency, and associate positive or negative sentiments. This method enhances product feedback analysis, helping businesses understand customer preferences

and improve services. The paper's novel approach to feature extraction and sentiment classification has influenced many subsequent studies in opinion mining.

Erik Cambria and Amir Hussain[3] presented a framework that integrates **common-sense knowledge** to enhance sentiment analysis accuracy. The Sentic Computing approach leverages concept-level semantics to analyze text sentiment more effectively than traditional word-level methods. By combining **NLP techniques with knowledge representation models**, Sentic Computing improves the understanding of emotions, intentions, and opinions in text. This work is pivotal in advancing sentiment analysis beyond surface-level language patterns.

Soujanya Poria et. al[4] reviewed the progress in affective computing, emphasizing the role of multimodal fusion. The authors discuss **integrating text, audio, and visual cues to improve sentiment analysis accuracy**. This multimodal approach addresses the limitations of unimodal systems and enhances the interpretation of emotional states. The comprehensive review highlights key challenges, including data integration techniques, feature extraction, and model fusion strategies.

Perna Chikersal, et. al[5] proposed a **hybrid sentiment analysis model** that integrates **rule-based methods with supervised learning techniques**. This combined approach improves accuracy by leveraging the strengths of both methods. The paper highlights techniques for feature engineering, data preprocessing, and model optimization. This innovative strategy effectively handles complex language patterns, making it valuable for analyzing social media data, customer reviews, and news articles.

Benjamin Snyder and Regina Barzilay[6] introduced the **Good Grief algorithm for ranking multiple aspects in sentiment analysis**. Their method assigns importance to different aspects of a product or service by analyzing review text. The algorithm efficiently **handles contradictory opinions within reviews and identifies dominant sentiment trends**. This approach proves valuable in improving recommendation systems and personalized marketing strategies.

Yan Qu, James Shanahan, and Janyce Wiebe[7] explored theoretical frameworks and practical applications of attitude and affect in textual data. Their work highlights the **psychological and linguistic aspects of emotion detection, emphasizing context-based sentiment interpretation**. The paper discusses innovative techniques for identifying sentiment trends in news articles, social media, and online communities, enhancing the understanding of human emotions in text.

Moshe Koppel and Jonathan Schler[8] focused on the significance of neutral examples in training sentiment analysis models. Their research reveals that incorporating neutral samples improves model robustness and reduces bias. The paper emphasizes the importance of neutral data in enhancing classification precision, particularly in scenarios involving ambiguous or emotionally complex content.

Filipe Nunes Ribeiro and Matheus Araujo[9] provided a comparative analysis of contemporary sentiment analysis techniques. The paper evaluates machine learning models, deep learning frameworks, and hybrid approaches based on performance metrics such as accuracy, precision, and recall. The authors identify strengths and weaknesses in each method, offering insights into optimal techniques for specific sentiment analysis tasks.

The study by Maite Taboada and Julian Brooke[10] presented a comprehensive overview of **lexicon-based sentiment analysis** methods. The authors delve into the use of manually curated and automatically generated lexicons to assign sentiment scores to words. The paper explores various approaches for combining individual word scores to evaluate text sentiment. The authors highlight challenges such as context dependency, negation handling, and ambiguity, proposing solutions like **sentiment shifters and discourse analysis**. The research emphasizes that lexicon-based methods perform well in scenarios with limited labeled data and provide interpretability in sentiment decisions. The study further demonstrates the application of lexicon-based approaches in social media analysis, review aggregation, and opinion mining. Overall, this paper provides valuable insights into the strengths and limitations of lexicon-based techniques, illustrating their role in both standalone analysis and ensemble models.

This study by Łukasz Augustyniak et al.[11] investigated the combination of **lexicon-based methods with ensemble learning for sentiment analysis**. The authors analyze multiple ensemble architectures, including bagging, boosting, and stacking, to improve the predictive performance of lexicon-based classifiers. The paper emphasizes the importance of feature engineering, integrating semantic and syntactic cues for enhanced text representation. By combining lexicon scores with machine learning classifiers, the authors demonstrate improved accuracy across various datasets, including movie reviews, tweets, and product feedback. The study also highlights ensemble techniques' robustness in handling noisy and imbalanced data. The proposed framework effectively addresses the limitations of standalone lexicon methods by leveraging diverse feature spaces. The authors provide detailed evaluation metrics, illustrating the potential of ensemble methods for enhancing sentiment analysis accuracy. This comprehensive study offers a valuable guide for researchers looking to implement lexicon-based ensemble classifiers in real-world applications.

The study by Mike Thelwall et al.[12] focused on developing a sentiment strength detection model tailored for **short informal texts**, such as social media posts, chat messages, and online reviews. The authors introduce the **SentiStrength tool**, which combines **a lexicon-based approach with specialized rules to assess sentiment polarity and intensity**. The study highlights the unique linguistic characteristics of informal text, including slang, abbreviations, and emoticons. SentiStrength effectively handles these complexities by integrating domain-specific sentiment lexicons and text normalization techniques. The tool's performance is evaluated across multiple datasets, demonstrating high accuracy in detecting both positive and

negative sentiment intensities. The authors emphasize that SentiStrength excels in real-time sentiment analysis, making it ideal for applications such as brand monitoring, customer feedback analysis, and public opinion tracking. This research provides a robust solution for assessing sentiment strength in brief and non-standardized text formats.

In this study, Bing Liu et al.[13] presented the Opinion Observer, a system designed to analyze and compare user opinions from web-based reviews. The paper introduces a novel aspect-based sentiment analysis framework that extracts product features and corresponding sentiments from text. The authors emphasize the importance of identifying opinion targets, employing natural language processing techniques such as **part-of-speech tagging and dependency parsing**. The Opinion Observer visualizes sentiment trends, enabling users to compare competing products based on customer feedback. The system's effectiveness is demonstrated in domains like consumer electronics, hotels, and restaurants. By focusing on feature-level sentiment analysis, the Opinion Observer provides granular insights into consumer preferences. The study highlights the value of combining sentiment analysis with data visualization for better decision-making in marketing, product development, and customer experience enhancement.

This study by Mario Cataldi et al.[14] proposed a novel approach for identifying feature-specific sentiments in user-generated reviews. The authors introduce a feature-sentiment extraction model that identifies key aspects such as **"location," "service," or "food" in reviews**. The model combines rule-based extraction with machine learning classifiers to detect sentiment polarity for each identified feature. The research emphasizes the importance of contextual understanding in opinion mining, addressing challenges such as **sarcasm, comparative language, and ambiguous expressions**. The authors evaluate their model on diverse review datasets, demonstrating high precision and recall. The study underscores the significance of feature-level sentiment analysis in applications like **hospitality, e-commerce, and product comparison platforms**. This research contributes practical insights into understanding nuanced consumer opinions, enhancing targeted marketing and product improvement strategies.

The study by Zhongwu Zhai et al. explored advancements in knowledge discovery techniques applied to sentiment analysis. The authors highlight the role of **data mining algorithms** in extracting meaningful insights from large-scale datasets. They introduce novel **clustering methods, feature selection techniques, and improved sentiment scoring models**. The research emphasizes leveraging knowledge graphs and semantic analysis for improved opinion mining. The authors demonstrate their approach through applications in product reviews, social media analysis, and customer satisfaction measurement. Their findings indicate that integrating data mining principles enhances sentiment prediction accuracy and **helps identify hidden patterns in textual data**. The study offers valuable insights for practitioners seeking to apply data mining strategies to sentiment analysis challenges.

Bin Liang's study introduced a **graph convolutional network (GCN)** framework enhanced with affective knowledge to improve aspect-based sentiment analysis. The authors propose incorporating sentiment-bearing words and affective knowledge into GCN models to capture complex semantic relationships. This approach significantly improves the model's ability to understand contextual sentiment cues. The paper demonstrates improved performance on benchmark datasets, particularly for distinguishing nuanced opinions on various product aspects. The research highlights the importance of integrating affective knowledge for tasks requiring fine-grained sentiment detection, particularly in customer reviews, social platforms, and user feedback systems.

Yukun Ma's[15] research explored **embedding commonsense knowledge into an attentive LSTM model** for targeted aspect-based sentiment analysis. The proposed method enhances traditional LSTM models by integrating knowledge graphs that provide background information on common phrases and expressions. This method improves the model's understanding of implicit sentiments and contextual clues. The authors evaluate their framework on review datasets, demonstrating enhanced accuracy in capturing opinion targets. The study emphasizes the importance of combining commonsense knowledge with neural network architectures to improve targeted sentiment detection for diverse domains.

Raksha Sharma et al.[16] presented a method for ranking sentiment intensity among adjectives using sentiment-bearing word embeddings. The authors introduce a novel ranking algorithm that leverages vector space representations of adjectives to predict their emotional intensity. The model efficiently distinguishes between strong, moderate, and weak adjectives based on contextual cues. Evaluation across multiple datasets reveals improved sentiment intensity prediction accuracy. This study offers insights into refining sentiment scoring methods and enhancing emotion detection frameworks in NLP applications.

M. S. Akhtar et al.[17] presented a **stacked ensemble approach** for predicting sentiment and emotion intensities. The proposed model combines **multiple classifiers to improve robustness** and accuracy. The authors employ diverse feature extraction techniques, including semantic embeddings, word-level cues, and syntactic features. By integrating these elements, the ensemble model effectively **predicts fine-grained sentiment intensity**. The study demonstrates the method's effectiveness in emotion analysis tasks, particularly in social media content and customer reviews. The research emphasizes ensemble methods' potential in enhancing emotion intensity prediction accuracy across diverse text sources.

II. METHODOLOGY

1. Introduction

This section outlines the methodology employed for text data preprocessing, sentiment encoding, data splitting, and the

application of BERT (Bidirectional Encoder Representations from Transformers) for text analysis. The methodology is divided into distinct stages to ensure clarity and reproducibility.

2. Data Preprocessing

The text preprocessing pipeline was designed to clean the input data by performing multiple text-cleaning steps. The following steps were applied:

2.1 Cleaning Pipeline

A custom function, $f(x)$, was implemented to standardize text data:

- **Lowercasing:**

$$f(x) = extlowercase(x)$$

- **Contraction Expansion:**

$$f(x) = extexpand_contractions(f(x))$$

- **Email Removal:**

$$f(x) = f(x) - extemails(f(x))$$

- **HTML Tag Removal:**

$$f(x) = f(x) - exthtml_tags(f(x))$$

- **Special Character Removal:**

$$f(x) = f(x) - extspecial_chars(f(x))$$

- **Accent Removal:**

$$f(x) = f(x) - extaccents(f(x))$$

The resulting text data after applying $f(x)$ is cleaner and ready for further processing.

3. Feature Engineering

3.1 Word Count Calculation

The number of words in each text entry was calculated using the following formula:

$$extnum_words(x) = | extSplit(x) |$$

Where $|\cdot|$ represents the cardinality of the split text array.

3.2 Sentiment Encoding

Categorical sentiment labels were encoded into numerical values using the following mapping:

extSentiment_Encoded
= {*extanger*: 0, *extfear*: 1, *extjoy*: 2, *extlove*: 3,
extsadness: 4, *extsurprise*: 5}

4. Data Splitting

The dataset was split into training and testing subsets using stratified sampling:

$D_{train}, D_{test} = extSplit(D, p = 0.7)$

Where:

- D is the original dataset.
- D_{train} and D_{test} are the resulting train and test subsets.
- p=0.7 indicates a 70:30 split ratio to ensure balance in sentiment distribution.

5. One-hot Encoding

To prepare sentiment labels for multi-class classification, categorical labels were converted into one-hot encoded vectors:

$y_{train} = extOneHot(D_{train}[extSentiment])$
 $y_{test} = extOneHot(D_{test}[extSentiment])$

6. Transformer Model Setup

A BERT model was chosen for its powerful text representation capabilities. The BERT tokenizer and model were defined as follows:

- **Tokenizer:**

$T = extBERT_Tokenizer(ext'bert - base - cased')$

- **Model:**

$M = extBERT_Model(ext'bert - base - cased')$

1. Tokenization Process

Let T be the input text sequence.

- The tokenizer maps the text sequence into a sequence of tokens:

$Tokens = Tokenizer(T)$

- Each token is then mapped to its corresponding integer ID using a vocabulary mapping function V:

$x = [V(t_1), V(t_2), \dots, V(t_n)]$

Where:

- $T = [t_1, t_2, \dots, t_n]$ is the sequence of tokens.
- $V: Token \rightarrow Z^+$ is the vocabulary mapping function.
- $x \in Z^n$ is the sequence of token IDs.

2. Embedding Layer

The token IDs are converted into dense vectors through an embedding matrix E:

$X = E \cdot x$

Where:

- $E \in R^{|V| \times d}$ is the embedding matrix.
- d is the embedding dimension.
- $X \in R^{n \times d}$ is the resulting sequence of word embeddings.

3. Positional Encoding

To encode positional information, a positional embedding matrix P is added:

$H_0 = X + P$

Where:

- $P \in R^{n \times d}$ is the positional encoding matrix.
- H_0 is the combined input representation.

4. Transformer Encoder Layers

Each transformer encoder layer applies self-attention and feed-forward layers. Let H_l be the output of the l-th encoder layer:

$$Hl + 1 = FFN(MultiHeadAttn(Hl))$$

Where:

- MultiHeadAttn is the multi-head self-attention mechanism.
- FFN is the position-wise feed-forward network.

1) Multi-Head Attention:

$$MultiHeadAttn(H) = Concat(head1, \dots, headh)WO$$

Each attention head is defined as:

$$head_i = softmax(Q_i K_i^T d_k) V_i$$

Where:

- $Q_i = H W_i Q$
- $K_i = H W_i K$
- $V_i = H W_i V$
- $W_i Q, W_i K, W_i V \in R_d \times d_k$ are learnable weight matrices.
- $d_k = d/h$ is the dimension of each attention head.

5. Final Output

The final output is a contextualized representation for each token:

$$Y = HL$$

Where:

- L is the total number of transformer layers.
- $Y \in R_n \times d$ is the final contextualized embedding.

6. Classification Head (Optional)

For downstream tasks like classification, a fully connected layer with softmax can be applied on the [CLS] token's representation:

$$\hat{y} = softmax(W \cdot Y[CLS] + b)$$

Where:


- W and b are learnable parameters.
- $Y[CLS]$ is the representation of the [CLS] token.

III. RESULTS

The following section presents the evaluation metrics for our sentiment analysis model. The classification_report output provides key performance indicators such as **Precision**, **Recall**, **F1-Score**, and **Support**, which are crucial in understanding the model's effectiveness.

1) Classification Report Analysis

Class	Precision	Recall	F1-Score	Support
0	0.91	0.96	0.94	813
1	0.94	0.85	0.89	712
2	0.92	0.97	0.95	2028
3	0.93	0.73	0.82	492
4	0.98	0.95	0.97	1739
5	0.72	0.97	0.83	216
Overall Metrics		Value		
Accuracy		0.93		
Macro Avg		Precision: 0.90 / Recall: 0.91 / F1-Score: 0.90		
Weighted Avg		Precision: 0.93 / Recall: 0.93 / F1-Score: 0.93		



2) Detailed Explanation

1. Class 0

- **Precision (0.91):** Out of all instances predicted as Class 0, 91% were correct.
- **Recall (0.96):** Out of all actual instances of Class 0, 96% were correctly identified.
- **F1-Score (0.94):** The high F1-score indicates strong overall performance for this class.

2. Class 1

- **Precision (0.94):** High precision shows the model effectively minimizes false positives for this class.
- **Recall (0.85):** Slightly lower recall suggests the model missed some actual Class 1 instances.
- **F1-Score (0.89):** A balanced score, indicating room for improvement in capturing more true positives.

3. Class 2

- **Precision (0.92):** Indicates excellent performance in correctly identifying Class 2.
- **Recall (0.97):** The model successfully identified 97% of actual Class 2 samples.
- **F1-Score (0.95):** Demonstrates robust performance with minimal errors.

4. Class 3




- **Precision (0.93):** The model correctly identifies Class 3 instances with high accuracy.
- **Recall (0.73):** A lower recall suggests more false negatives in this category.
- **F1-Score (0.82):** This score reflects the trade-off between precision and recall for this class.

5. Class 4

- **Precision (0.98):** Outstanding precision demonstrates excellent prediction accuracy for this class.
- **Recall (0.95):** The model successfully identified 95% of actual Class 4 instances.
- **F1-Score (0.97):** Excellent overall performance with minimal errors.

6. Class 5

- **Precision (0.72):** Precision is lower, indicating some misclassifications.
- **Recall (0.97):** High recall shows the model successfully identified almost all Class 5 instances.
- **F1-Score (0.83):** Despite lower precision, the strong recall contributes to a solid F1 score.



Overall Accuracy: 0.9300
Overall Precision (Macro Avg): 0.9018
Overall Recall (Macro Avg): 0.9053
Overall F1-Score (Macro Avg): 0.8977

3) Overall Performance

- **Accuracy (0.93):** The model correctly predicted 93% of the total test samples, indicating high reliability.
- **Macro Average:** Since this averages all class metrics equally, it highlights overall consistency.
- **Weighted Average:** Given this metric weights classes based on sample size, it emphasizes the model's effectiveness across both majority and minority classes.

4) Evaluation Metrics Explained

1. Accuracy

$$\text{Accuracy} = \frac{\text{Correct Predictions}}{\text{Total Predictions}}$$

- Accuracy indicates the overall percentage of correct predictions across all classes.

2. Macro Average Precision

$$\text{Macro Avg Precision}$$

$$= \frac{\text{Precision}_0 + \text{Precision}_1 + \dots + \text{Precision}_n}{n}$$

- Precision measures the proportion of true positives among predicted positive samples.
- The macro average gives equal importance to all classes, making it suitable when class imbalance exists.

3. Macro Average Recall

$$\text{Macro Avg Recall} = \frac{\text{Recall}_0 + \text{Recall}_1 + \dots + \text{Recall}_n}{n}$$

- Recall reflects the model's ability to correctly identify actual positive samples.
- The macro average recall ensures minority classes are treated equally.

4. Macro Average F1-Score

$$\text{F1-Score} = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

- The F1-score is a harmonic mean of precision and recall, balancing the trade-off between false positives and false negatives.
- The macro average F1-score combines all class F1-scores equally.

5) Conclusion

The model demonstrates robust performance across most classes, with particularly strong results in Classes 0, 2, and 4. However, Class 3 and Class 5 show lower recall and precision, indicating potential areas for improvement. Future enhancements may involve techniques such as **data augmentation, class balancing, or hyperparameter tuning** to improve performance in these underperforming classes. The high **accuracy** and **macro average scores** demonstrate the model's robustness in classifying sentiments across multiple classes. However, if performance discrepancies exist between minority and majority classes, strategies like **oversampling, undersampling, or class-weight adjustments** can further enhance the model's effectiveness.

REFERENCES

- [1] B. Pang and L. Lee, "Opinion mining and sentiment analysis," *Foundations and Trends in Information Retrieval*, vol. 2, no. 1–2, pp. 1–135, 2008.
- [2] M. Hu and B. Liu, "Mining and summarizing customer reviews," in *Proceedings of the 10th ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD)*, Seattle, WA, USA, 2004, pp. 168–177.
- [3] E. Cambria and A. Hussain, *Sentic Computing: A Common-Sense-Based Framework for Concept-Level Sentiment Analysis*. Cham, Switzerland: Springer, 2012.
- [4] S. Poria, E. Cambria, R. Bajpai, and A. Hussain, "A review of affective computing: From unimodal analysis to multimodal fusion," *Information Fusion*, vol. 37, pp. 98–125, Sept. 2017.
- [5] P. Chikersal, M. Belgrave, and S. Wu, "Hybrid sentiment analysis: Leveraging rule-based and supervised learning approaches," in *Proceedings of the International Conference on Artificial Intelligence and Soft Computing (ICAISC)*, Zakopane, Poland, 2019, pp. 153–164.
- [6] B. Snyder and R. Barzilay, "Multiple aspect ranking using the Good Grief algorithm," in *Proceedings of the 43rd Annual Meeting of the Association for Computational Linguistics (ACL)*, Ann Arbor, MI, USA, 2005, pp. 161–168.
- [7] Y. Qu, J. Shanahan, and J. Wiebe, "Attitude and affect in text: Theoretical foundations and practical applications," in *Handbook of Natural Language Processing*, 2nd ed., Boca Raton, FL, USA: CRC Press, 2010, pp. 677–700.
- [8] M. Koppel and J. Schler, "Neutral examples improve text sentiment classification," in *Proceedings of the 19th International Conference on Computational Linguistics (COLING)*, Taipei, Taiwan, 2002, pp. 417–423.
- [9] F. N. Ribeiro and M. Araujo, "A comparative study of deep learning techniques for sentiment analysis," *IEEE Transactions on Affective Computing*, vol. 12, no. 4, pp. 1015–1028, Oct.–Dec. 2021.
- [10] M. Taboada and J. Brooke, "Lexicon-based methods for sentiment analysis," *Computational Linguistics*, vol. 38, no. 1, pp. 267–307, 2012.
- [11] Ł. Augustyniak, R. Rabciga-Wiśniewska, and P. Szymczak, "Ensemble learning for sentiment analysis using lexicon-based approaches," in *Proceedings of the 2020 IEEE International Conference on Artificial Intelligence and Data Science (ICAIDS)*, Delhi, India, 2020, pp. 1–6.
- [12] M. Thelwall, K. Buckley, and G. Paltoglou, "SentiStrength: A sentiment strength detection tool for social web data," in *Proceedings of the 2010 International Conference on Web Intelligence and Intelligent Agent*

Technology (WI-IAT), Toronto, Canada, 2010, pp. 119–124.

- [13] B. Liu, M. Hu, and J. Cheng, “Opinion Observer: Analyzing and comparing opinions on the web,” in *Proceedings of the 14th International World Wide Web Conference (WWW)*, Chiba, Japan, 2005, pp. 342–351.
- [14] M. Cataldi, L. Di Caro, and C. Schifanella, “Feature-specific sentiment analysis for user-generated reviews,” *ACM Transactions on Information Systems*, vol. 36, no. 3, pp. 1–28, May 2018.
- [15] Y. Ma and E. Cambria, “Embedding commonsense knowledge into an attentive LSTM model for aspect-based sentiment analysis,” in *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, Hong Kong, China, 2019, pp. 302–310.
- [16] R. Sharma, P. Kumar, and S. Singh, “Sentiment intensity ranking using word embeddings,” in *Proceedings of the 2018 IEEE International Conference on Computational Intelligence and Knowledge Economy (ICCIKE)*, Dubai, UAE, 2018, pp. 223–228.
- [17] M. S. Akhtar, A. Ekbal, and P. Bhattacharyya, “A stacked ensemble framework for sentiment and emotion intensity prediction,” *Information Processing & Management*, vol. 57, no. 4, pp. 102–113, 2020.